

RESPONSIBLE AI ADOPTION TOOLKIT FOR SMES â V0.2 PUBLIC PACKET

Responsible AI at Work Lab

Project page: <https://murphylawtex-bit.github.io/responsible-ai-at-work-lab/>

Repository: <https://github.com/murphylawtex-bit/responsible-ai-at-work-lab>

Contact: murphy.law.tex@gmail.com

QUICKSTART: RUN A RESPONSIBLE AI PILOT REVIEW

STEP 1 â DEFINE THE USE CASE

Complete 'templates/01-ai-use-case-intake.md'.

If the use case is vague, stop and narrow it.

STEP 2 â MAP AFFECTED PEOPLE

Complete 'templates/02-stakeholder-and-impact-map.md' and 'templates/04-workforce-impact-checklist.md'.

If affected workers or users cannot provide feedback, do not scale the pilot.

STEP 3 â MEASURE BEFORE CHANGING THE WORKFLOW

Complete 'templates/03-productivity-roi-measurement-plan.md'.

If there is no baseline, productivity claims are guesses.

STEP 4 â REGISTER RISKS

Complete 'templates/05-risk-register.md'.

Every high-impact risk needs an owner, mitigation, and review date.

STEP 5 â DEFINE HUMAN OVERSIGHT

Complete 'templates/06-human-oversight-plan.md'.

If nobody is accountable when AI is wrong, the pilot is not ready.

STEP 6 â RUN A LIMITED PILOT

Use a narrow workflow, limited data, and clear stop conditions.

STEP 7 â REVIEW BEFORE SCALE

Complete 'templates/08-pilot-review-scorecard.md'.

Decision options:

- â Stop
- â Revise and rerun
- â Continue limited pilot
- â Expand carefully
- â Scale with controls

ONE-PAGE QUICKSTART

RESPONSIBLE AI ADOPTION REVIEW

Use this before piloting AI in a business workflow.

1. DEFINE THE USE CASE

What workflow changes? What AI does what? Who owns it?

Template: '01-ai-use-case-intake.md'

2. IDENTIFY AFFECTED PEOPLE

Who uses it? Who is affected by it? Could it change job quality, access, fairness, privacy, or trust?

Templates: '02-stakeholder-and-impact-map.md', '04-workforce-impact-checklist.md'

3. MEASURE VALUE

What is the baseline? What improvement is expected? How will quality and rework be measured?

Template: '03-productivity-roi-measurement-plan.md'

4. REGISTER RISK

What can go wrong? Who owns the risk? What controls are required?

Template: '05-risk-register.md'

5. KEEP HUMANS ACCOUNTABLE

What decisions stay human-owned? Who reviews AI output? Who can override it?

Template: '06-human-oversight-plan.md'

6. PILOT NARROWLY

Use limited scope, limited data, and clear stop conditions.

Guide: 'implementation-guide.md'

7. REVIEW BEFORE SCALE

Stop, revise, continue, expand, or scale only after evidence review.

Template: '08-pilot-review-scorecard.md'

30-DAY RESPONSIBLE AI ADOPTION PILOT GUIDE

OBJECTIVE

Run a limited, measurable AI adoption pilot that creates real business evidence while protecting workers, customers, and public trust.

WEEK 1 â DEFINE AND SCREEN

1. Pick one narrow workflow.
2. Complete '01-ai-use-case-intake.md'.
3. Complete '02-stakeholder-and-impact-map.md'.
4. Create initial '05-risk-register.md'.
5. Decide whether the use case is safe enough for a limited pilot.

Deliverable: pilot charter and risk screen.

WEEK 2 â MEASURE AND GOVERN

1. Complete '03-productivity-roi-measurement-plan.md'.
2. Complete '06-human-oversight-plan.md'.
3. Adopt or adapt '07-ai-governance-policy.md'.

4. Brief affected workers/users.

Deliverable: approved pilot plan.

WEEKS 3 & 4 â PILOT

1. Run the AI workflow in parallel or limited production.
2. Track time, quality, rework, review burden, and incidents.
3. Collect worker/user feedback.
4. Update the risk register weekly.

Deliverable: pilot evidence file.

END OF MONTH â REVIEW

1. Complete '08-pilot-review-scorecard.md'.
2. Decide: stop, revise, expand, or scale.
3. Document conditions required for next stage.

Deliverable: adoption decision memo.

GOOD PILOT CHARACTERISTICS

- â Narrow enough to reverse.
- â Clear human owner.
- â Baseline exists.
- â No high-impact decisions without special review.
- â Workers/users can give feedback.
- â Risks are tracked visibly.

BAD PILOT CHARACTERISTICS

- â Vague productivity claims.
- â No baseline.
- â No human accountability.
- â Sensitive data in unapproved tools.
- â Hidden worker surveillance.
- â â Human in the loopâ that is really rubber-stamping.

NIST AI RISK MANAGEMENT FRAMEWORK CROSSWALK

PURPOSE

This crosswalk maps the Responsible AI Adoption Toolkit for SMEs to the NIST AI Risk Management Framework (AI RMF) functions: **Govern, Map, Measure, and Manage**.

The goal is not to turn small organizations into compliance departments. The goal is to translate responsible AI risk management into practical workflow-level questions and documents that SMEs can actually use.

TOOLKIT-TO-NIST SUMMARY

NIST AI RMF Function	SME Toolkit Component	Practical purpose
Govern	AI Governance Policy	Defines principles, ownership, prohibited uses, review cadence, and incident response.
Govern	Human Oversight Plan	Assigns human accountability, override rights, review requirements, and stop conditions.
Govern	Pilot Review Scorecard	Creates accountable scale/stop/revise decisions.
Map	AI Use Case Intake	Defines context, purpose, users, data, affected people, and AI role.
Map	Stakeholder and Impact Map	Identifies affected groups, likely benefits, harms, and engagement needs.
Map	Workforce Impact Checklist	Maps job-quality, autonomy, surveillance, skill, and role-change implications.
Measure	Productivity and ROI Measurement Plan	Establishes baseline, metrics, target outcomes, and evidence quality.
Measure	Risk Register	Measures likelihood, impact, controls, ownership, and status.
Measure	Pilot Review Scorecard	Compares results against baseline, including risk and worker experience.
Manage	Risk Register	Tracks mitigations, owners, dates, and open/closed status.
Manage	Human Oversight Plan	Defines operational controls, escalation, review, and appeal paths.
Manage	Implementation Guide	Provides staged pilot workflow and review process.

GOVERN

NIST emphasis: establish policies, processes, accountability, organizational roles, and risk culture.

Toolkit support:

• '07-ai-governance-policy.md'
• '06-human-oversight-plan.md'
• '08-pilot-review-scorecard.md'

SME questions:

1. Who owns this AI use case?
2. Who is accountable when the AI is wrong?
3. What uses are prohibited without special review?
4. What decisions must remain human-owned?
5. Who can pause or stop the pilot?
6. How often is the AI workflow reviewed?
7. How are incidents logged and escalated?

Minimum evidence:

- â€¢ Named sponsor and process owner
- â€¢ Human review requirements
- â€¢ Stop conditions
- â€¢ Incident response path
- â€¢ Pilot review date

MAP

NIST emphasis: understand context, intended purpose, affected groups, assumptions, benefits, and risks.

Toolkit support:

- â€¢ '01-ai-use-case-intake.md'
- â€¢ '02-stakeholder-and-impact-map.md'
- â€¢ '04-workforce-impact-checklist.md'

SME questions:

1. What exact workflow is changing?
2. What AI capability is being used?
3. What data goes in and what output comes out?
4. Who directly uses the system?
5. Who is affected indirectly?
6. Could this change job quality, autonomy, pay, scheduling, monitoring, or advancement?
7. Could this affect customers, applicants, vendors, or public access?

Minimum evidence:

- â€¢ Use-case description
- â€¢ Data/input/output description
- â€¢ Stakeholder list
- â€¢ Workforce impact notes
- â€¢ Initial risk screen

MEASURE

NIST emphasis: analyze, evaluate, benchmark, and monitor AI risks and performance.

Toolkit support:

- â€¢ '03-productivity-roi-measurement-plan.md'
- â€¢ '05-risk-register.md'
- â€¢ '08-pilot-review-scorecard.md'

SME questions:

1. What is the current baseline?
2. What productivity gain is expected?
3. How will quality be measured?
4. How will errors/rework be tracked?
5. How will worker experience be measured?
6. What risks are high-impact even if unlikely?

7. What evidence would prove the pilot failed?

Minimum evidence:

- â€ Baseline metrics
- â€ Success thresholds
- â€ Risk likelihood/impact ratings
- â€ Incident log
- â€ Worker/user feedback

MANAGE

NIST emphasis: prioritize, respond to, treat, monitor, and communicate AI risks.

Toolkit support:

- â€ '05-risk-register.md'
- â€ '06-human-oversight-plan.md'
- â€ 'docs/implementation-guide.md'
- â€ '08-pilot-review-scorecard.md'

SME questions:

1. Which risks must be mitigated before pilot?
2. Which risks can be monitored during pilot?
3. Who owns each mitigation?
4. What triggers pause/stop/escalation?
5. What must change before scale-up?
6. What should be communicated to workers or affected people?

Minimum evidence:

- â€ Risk owners
- â€ Mitigation actions
- â€ Review dates
- â€ Stop/revise/scale decision
- â€ Conditions for expansion

LIGHTWEIGHT SCORING MODEL

Use a 1â 5 score for each dimension:

- | Dimension | 1 = weak/high risk | 5 = strong/low risk |
- |---|---|---|
- | Purpose clarity | Vague or speculative | Narrow, useful, measurable |
- | Data sensitivity | Sensitive/uncontrolled | Minimal/approved data |
- | Human accountability | Unclear | Named owner and review path |
- | Workforce impact | Unknown/possibly harmful | Reviewed with affected workers |
- | Measurement quality | No baseline | Baseline and success metrics defined |
- | Risk controls | No mitigations | Controls, owners, stop rules defined |
- | Reversibility | Hard to undo | Limited pilot, easy rollback |

Suggested decision rule:

â€¢ **28â 35:** Pilot likely appropriate with normal controls.
â€¢ **20â 27:** Revise before pilot; address weak dimensions.
â€¢ **Below 20:** Do not pilot until governance and risk gaps are fixed.
â€¢ Any high-impact use case involving employment, pay, safety, legal, financial, health, housing, education, or access rights requires special review regardless of score.

GRANT POSITIONING NOTE

This crosswalk strengthens funding applications by showing that the toolkit is not merely a checklist collection. It is structured around a recognized AI risk-management framework while remaining usable for SMEs.

AI USE CASE INTAKE

1. BASIC INFORMATION

â€¢ Organization / team:
â€¢ Date:
â€¢ Sponsor / owner:
â€¢ Process owner:
â€¢ Proposed AI use case name:
â€¢ Pilot or production?:

2. PROBLEM STATEMENT

What business problem or workflow friction are we trying to solve?

3. PROPOSED AI-ENABLED WORKFLOW

Describe what the AI will do and where it fits in the process.

4. CURRENT PROCESS BASELINE

â€¢ Current process owner:
â€¢ Current tools:
â€¢ Current average cycle time:
â€¢ Current error/rework rate:
â€¢ Current cost estimate:
â€¢ Current pain points:

5. INTENDED OUTCOMES

Select all that apply:

- ☐ Save time
- ☐ Improve quality
- ☐ Reduce rework/errors
- ☐ Improve customer/user experience
- ☐ Improve access or inclusion
- ☐ Improve documentation/knowledge capture
- ☐ Support decision-making
- ☐ Reduce risk
- ☐ Other:

6. USERS AND AFFECTED PARTIES

- ☐ Direct users:
- ☐ Managers/approvers:
- ☐ Internal affected teams:
- ☐ External affected parties/customers:
- ☐ Workers whose roles may change:

7. AI SYSTEM DETAILS

- ☐ AI tool/model/vendor:
- ☐ Data inputs:
- ☐ Data outputs:
- ☐ Does it use personal/sensitive data?:
- ☐ Does it make recommendations or decisions?:
- ☐ Does it interact with external people?:

8. HUMAN ACCOUNTABILITY

- ☐ Who reviews AI output?
- ☐ Who can override it?
- ☐ Who is accountable if it is wrong?
- ☐ What decisions are AI prohibited from making?

9. INITIAL RISK SCREEN

- ☐ Could the AI output materially affect someone's job, pay, access, safety, reputation, or rights?
- ☐ Could confidential or personal data be exposed?
- ☐ Could bias or unfair treatment occur?
- ☐ Could hallucinations create operational or legal risk?
- ☐ Could employees feel surveilled, deskilled, or displaced?

10. GO/NO-GO FOR DEEPER REVIEW

- â€¢ ☐ Proceed to impact map
- â€¢ ☐ Needs revision before review
- â€¢ ☐ Do not pursue

Reason:

STAKEHOLDER AND IMPACT MAP

USE CASE

â€¢ Use case name:

â€¢ Date:

â€¢ Prepared by:

STAKEHOLDER TABLE

Stakeholder group	How they interact with the workflow	Potential benefit	Potential harm/risk	Engagement needed?	Notes
---	---	---	---	---	---
Employees/users		Yes/No			
Managers		Yes/No			
Customers/public		Yes/No			
IT/security		Yes/No			
Compliance/legal		Yes/No			
Other		Yes/No			

WORKFORCE IMPACT QUESTIONS

1. Which tasks become automated, accelerated, or changed?
2. Which tasks remain human-owned?
3. Will the AI increase worker autonomy or reduce it?
4. Will workers need new skills or training?
5. Could the system be used for monitoring or performance discipline?
6. Could the system reduce headcount, hours, wages, or advancement opportunity?
7. How will affected workers provide feedback before scale-up?

SOCIETAL IMPACT QUESTIONS

1. Does this use case change access to services, information, jobs, or opportunities?
2. Could it create unequal outcomes across groups?
3. Does it increase dependence on opaque systems?
4. Does it improve or reduce transparency?
5. What would responsible public disclosure look like?

ENGAGEMENT PLAN

- â€¢ Stakeholders to interview:
- â€¢ Questions to ask:
- â€¢ Feedback deadline:
- â€¢ Changes made based on feedback:

PRODUCTIVITY AND ROI MEASUREMENT PLAN

USE CASE

- â€¢ Use case name:
- â€¢ Pilot period:
- â€¢ Owner:

BASELINE METRICS

Metric	Current baseline	Source	Confidence	Notes
Cycle time			Low/Med/High	
Labor hours per unit			Low/Med/High	
Error/rework rate			Low/Med/High	
Customer/user satisfaction			Low/Med/High	
Throughput			Low/Med/High	
Cost per unit			Low/Med/High	

PILOT SUCCESS METRICS

Metric	Target	Minimum acceptable	How measured	Review date
Time saved				
Quality improvement				
Rework reduction				
User satisfaction				
Risk incidents				

COST MODEL

- â€¢ Tool/vendor cost:
- â€¢ Setup/configuration time:
- â€¢ Training time:
- â€¢ Review/oversight time:
- â€¢ Maintenance cost:
- â€¢ Expected savings/value:

MEASUREMENT CAUTIONS

- â€ Do not count time saved if it merely shifts review burden elsewhere.
- â€ Separate novelty effect from durable productivity gain.
- â€ Track quality and rework, not only speed.
- â€ Track worker experience and customer impact alongside ROI.

DECISION RULE

- â€ Scale if:
- â€ Revise if:
- â€ Stop if:

WORKFORCE IMPACT CHECKLIST

USE CASE

- â€ Use case name:
- â€ Date:
- â€ Review owner:

JOB AND TASK IMPACT

- â€ ☐ Tasks automated or partially automated are documented.
- â€ ☐ Tasks requiring human judgment are documented.
- â€ ☐ New skills/training requirements are identified.
- â€ ☐ Potential role redesign is described.
- â€ ☐ Potential headcount/hour/pay impacts are disclosed internally before scale.

WORKER PARTICIPATION

- â€ ☐ Affected workers have reviewed the proposed workflow.
- â€ ☐ Feedback channels are available before and during pilot.
- â€ ☐ Workers can report harmful, inaccurate, or burdensome AI behavior.
- â€ ☐ Pilot results include worker feedback, not only management metrics.

JOB QUALITY

Assess expected impact:

Dimension	Improves	Neutral	Worsens	Unknown	Notes
Autonomy					
Skill development					
Workload intensity					
Monitoring/surveillance					
Fairness					
Safety/stress					

| Advancement opportunity | | | | |

GUARDRAILS

- ☐ AI outputs are not used as sole basis for hiring, firing, discipline, pay, promotion, or scheduling decisions.
- ☐ Workers know when AI is being used in the workflow.
- ☐ Workers know how to challenge or correct AI output.
- ☐ Human review is meaningful, not rubber-stamping.
- ☐ Training is provided before expectations change.

SUMMARY JUDGMENT

- ☐ Overall workforce impact: Positive / Mixed / Negative / Unknown
- ☐ Conditions required before pilot:
- ☐ Conditions required before scale:

AI RISK REGISTER

USE CASE

- ☐ Use case name:
- ☐ Owner:
- ☐ Last updated:

ID	Risk	Category	Likelihood	Impact	Controls/Mitigation	Owner	Status	Review date
---	---	---	---	---	---	---	---	---
R-001	Hallucinated or incorrect output	Quality	Med	High	Human review; source citation; sampling audit		Open	
R-002	Confidential data leakage	Privacy/Security	Low/Med/High	High	Data minimization; approved tools; no sensitive inputs		Open	
R-003	Bias or unfair treatment	Fairness	Low/Med/High	High	Impact testing; appeal path; human oversight		Open	
R-004	Worker deskilling or autonomy loss	Workforce	Low/Med/High	Med	Training; role redesign; worker feedback		Open	
R-005	Overreliance on AI output	Governance	Med	Med/High	Clear accountability; review checklist		Open	
R-006	Vendor/model change affects output	Operational	Med	Med	Version tracking; periodic revalidation		Open	

CATEGORIES

- ☐ Privacy/security
- ☐ Fairness/bias
- ☐ Quality/reliability
- ☐ Safety
- ☐ Legal/compliance

- â€ Workforce/job quality
- â€ Transparency/accountability
- â€ Operational resilience
- â€ Reputational/public trust

ESCALATION RULE

Pause the pilot if any high-impact risk lacks an owner, mitigation, or review process.

HUMAN OVERSIGHT PLAN

USE CASE

- â€ Use case name:
- â€ Owner:
- â€ Date:

AI ROLE

The AI system will:

- â€ ☐ Draft
- â€ ☐ Summarize
- â€ ☐ Classify
- â€ ☐ Recommend
- â€ ☐ Prioritize
- â€ ☐ Generate code/content
- â€ ☐ Analyze data
- â€ ☐ Interact with external users
- â€ ☐ Other:

HUMAN-OWNED DECISIONS

List decisions the AI may not make autonomously:

- 1.
- 2.
- 3.

REVIEW PROCESS

- â€ Who reviews AI output?
- â€ What percentage of outputs are reviewed?
- â€ What outputs require 100% review?
- â€ What evidence/source checks are required?
- â€ How are corrections logged?

OVERRIDE AND APPEAL

- â€¢ Who can override AI output?
- â€¢ How can affected people challenge output?
- â€¢ Who reviews disputes?
- â€¢ What is the expected response time?

ACCOUNTABILITY

- â€¢ Accountable executive/sponsor:
- â€¢ Process owner:
- â€¢ Technical owner:
- â€¢ Risk/compliance owner:
- â€¢ Workforce/stakeholder owner:

STOP CONDITIONS

Stop or pause the AI workflow if:

- â€¢ ☐ It produces repeated material errors.
- â€¢ ☐ It exposes confidential/personal data.
- â€¢ ☐ It creates unfair or discriminatory outcomes.
- â€¢ ☐ Workers/customers cannot understand or challenge its use.
- â€¢ ☐ Review burden exceeds productivity gains.
- â€¢ ☐ It is used outside approved scope.

LIGHTWEIGHT RESPONSIBLE AI GOVERNANCE POLICY

PURPOSE

This policy defines how the organization pilots and adopts AI tools responsibly, with attention to business value, human accountability, workforce impact, privacy, fairness, and public trust.

SCOPE

Applies to AI tools used for drafting, summarization, analysis, recommendations, automation, customer interaction, operational decision support, or employee-facing workflows.

PRINCIPLES

1. ****Human accountability**** â€” A person remains accountable for AI-assisted decisions.

2. **Purpose limitation** â AI is used only for approved purposes.
3. **Transparency** â Affected users/workers should know when AI materially shapes a workflow.
4. **Data minimization** â Use only the data needed; avoid sensitive/confidential inputs unless approved.
5. **Fairness and access** â Assess whether outcomes differ unfairly across groups.
6. **Worker participation** â Affected workers should provide feedback before scale-up.
7. **Measurement** â Productivity claims must be measured against a baseline.
8. **Reversibility** â Pilots should be limited and stoppable before full scale.

REQUIRED REVIEW BEFORE PILOT

- â AI Use Case Intake completed
- â Stakeholder and Impact Map completed
- â Risk Register created
- â Human Oversight Plan approved
- â Measurement Plan approved

PROHIBITED USES WITHOUT SPECIAL REVIEW

- â Sole basis for hiring, firing, discipline, pay, promotion, scheduling, lending, housing, insurance, legal, medical, safety, or other high-impact decisions
- â Covert employee surveillance
- â Uploading confidential, personal, regulated, or customer data into unapproved tools
- â External communications without human review
- â Automated impersonation of a person without disclosure and approval

PILOT REQUIREMENTS

- â Defined owner
- â Defined duration
- â Baseline metrics
- â Worker/user feedback channel
- â Incident log
- â Review meeting before expansion

REVIEW CADENCE

- â Pilot review: end of pilot or 30 days, whichever comes first
- â Production review: quarterly or after major model/vendor/workflow changes

INCIDENT RESPONSE

Material errors, data exposure, unfair outcomes, or scope violations must be logged, escalated to the owner, and reviewed before the workflow continues.

AI PILOT REVIEW SCORECARD

PILOT SUMMARY

- â€ Use case name:
- â€ Pilot dates:
- â€ Owner:
- â€ Participants:
- â€ Tools/models used:

RESULTS

Dimension	Target	Actual	Rating	Notes
Time saved			Red/Yellow/Green	
Quality			Red/Yellow/Green	
Rework/errors			Red/Yellow/Green	
User satisfaction			Red/Yellow/Green	
Worker experience			Red/Yellow/Green	
Risk incidents	0		Red/Yellow/Green	
Review burden			Red/Yellow/Green	

WHAT WORKED

-
-
-

WHAT FAILED OR CREATED FRICTION

-
-
-

WORKFORCE AND SOCIETAL IMPACT NOTES

- â€ Job quality impact:
- â€ Training needs:
- â€ Transparency concerns:
- â€ Fairness/access concerns:
- â€ Stakeholder feedback:

RISK REVIEW

- â€ New risks discovered:

â€¢ Controls that worked:
â€¢ Controls that failed:
â€¢ Required mitigations before expansion:

DECISION

â€¢ ☐ Stop
â€¢ ☐ Continue limited pilot
â€¢ ☐ Revise and rerun pilot
â€¢ ☐ Expand to another team/workflow
â€¢ ☐ Scale to production with controls

CONDITIONS FOR NEXT STAGE

- 1.
- 2.
- 3.

AI USE CASE READINESS AND RISK SCORING WORKSHEET

USE CASE

â€¢ Use case name:
â€¢ Date:
â€¢ Reviewer:

SCORING

Score each dimension from 1 to 5.

Dimension	1 = weak/high risk	3 = partial	5 = strong/low risk	Score	Notes
Purpose clarity	Vague, speculative	Useful but broad	Narrow, useful, measurable		
Data sensitivity	Sensitive/uncontrolled	Mixed/unclear	Minimal/approved data		
Human accountability	Unclear	Partially assigned	Named owner and review path		
Workforce impact	Unknown or likely harmful	Mixed/uncertain	Reviewed with affected workers		
Measurement quality	No baseline	Partial baseline	Baseline and success metrics defined		
Risk controls	No mitigations	Some mitigations	Controls, owners, stop rules defined		
Reversibility	Hard to undo	Partly reversible	Limited pilot, easy rollback		

TOTAL SCORE

Total: ____ / 35

DECISION GUIDANCE

â€¢ **28â 35:** Pilot likely appropriate with normal controls.
â€¢ **20â 27:** Revise before pilot; address weak dimensions.
â€¢ **Below 20:** Do not pilot until governance and risk gaps are fixed.

OVERRIDE RULES

Regardless of score, require special review if the use case affects:

â€¢ hiring, firing, discipline, pay, promotion, or scheduling;
â€¢ health, safety, legal, financial, lending, housing, insurance, or education outcomes;
â€¢ access to essential services;
â€¢ surveillance, monitoring, or automated evaluation of people;
â€¢ children, vulnerable populations, or protected classes.

REQUIRED ACTIONS BEFORE PILOT

- 1.
- 2.
- 3.

REQUIRED ACTIONS BEFORE SCALE-UP

- 1.
- 2.
- 3.

AI PILOT DECISION MEMO

USE CASE

â€¢ Use case name:
â€¢ Pilot dates:
â€¢ Decision date:
â€¢ Decision owner:

DECISION

- â€¢ [] Stop
- â€¢ [] Revise and rerun pilot
- â€¢ [] Continue limited pilot
- â€¢ [] Expand to another workflow/team
- â€¢ [] Scale with controls

SUMMARY

Briefly summarize the pilot and recommendation.

EVIDENCE REVIEWED

- â€¢ [] Use-case intake
- â€¢ [] Stakeholder/impact map
- â€¢ [] Productivity/ROI measurement plan
- â€¢ [] Workforce impact checklist
- â€¢ [] Risk register
- â€¢ [] Human oversight plan
- â€¢ [] Worker/stakeholder feedback
- â€¢ [] Pilot review scorecard
- â€¢ [] Incident log

RESULTS

Area	Finding	Confidence	Notes
Time/productivity		Low/Med/High	
Quality/rework		Low/Med/High	
Review burden		Low/Med/High	
Worker experience		Low/Med/High	
Customer/user impact		Low/Med/High	
Privacy/security		Low/Med/High	
Fairness/access		Low/Med/High	

CONDITIONS FOR APPROVAL

The next stage is approved only if these conditions are met:

- 1.
- 2.
- 3.

RISKS ACCEPTED

Risk	Why accepted	Owner	Review date

| | | |

RISKS NOT ACCEPTED

Risk	Required mitigation	Owner	Deadline
---	---	---	---

COMMUNICATION PLAN

â€ Who must be informed?
â€ What will be disclosed?
â€ How will feedback continue?

NEXT REVIEW DATE

WORKER AND STAKEHOLDER FEEDBACK GUIDE

PURPOSE

AI adoption changes work. This guide helps organizations collect practical feedback from people affected by an AI pilot before deciding whether to scale.

WHEN TO USE

Use this guide:

1. Before a pilot starts.
2. Mid-pilot if the workflow changes materially.
3. At pilot review before scale-up.

PRINCIPLES

â€ Feedback should be safe to give honestly.
â€ Workers should know what AI is being tested and why.
â€ Feedback should influence the pilot decision.
â€ Productivity metrics should not be the only evidence considered.

PRE-PILOT QUESTIONS

1. Which parts of this workflow are most frustrating or repetitive today?
2. Where do errors, delays, or rework usually happen?
3. What would make AI assistance useful here?
4. What would make AI assistance harmful or burdensome?
5. What information should not be put into an AI system?
6. Which decisions should remain human-owned?
7. What would you need to trust this workflow?

MID-PILOT QUESTIONS

1. Is the AI saving time, shifting work, or creating extra review burden?
2. Are outputs accurate enough to be useful?
3. Are people over-trusting the AI?
4. Are people ignoring the AI because it is not useful?
5. Has the pilot changed autonomy, workload, stress, or skill use?
6. Have any privacy, fairness, or transparency concerns appeared?
7. What should be changed before continuing?

END-OF-PILOT QUESTIONS

1. Should this workflow stop, revise, expand, or scale?
2. What evidence supports that decision?
3. What risks remain unresolved?
4. What training would be needed before wider use?
5. What should managers not do with this AI output?
6. What would make the workflow more transparent or fair?

FEEDBACK SUMMARY TEMPLATE

Question area	Main feedback	Evidence/examples	Required change	Owner
Time/workload				
Quality/rework				
Autonomy/job quality				
Privacy/security				
Fairness/access				
Trust/transparency				

MINIMUM STANDARD BEFORE SCALE-UP

Do not scale if:

- affected people were not informed;
- feedback was not collected;
- high-impact concerns remain unresolved;
- review burden exceeds productivity gain;
- human accountability is unclear.

EXAMPLE: CUSTOMER SUPPORT TRIAGE PILOT

SCENARIO

A 45-person company wants to use AI to categorize inbound support emails, summarize customer issues, suggest priority level, and draft internal routing notes.

AI ROLE

The AI classifies and summarizes. It does not send customer replies, close tickets, promise refunds, or make account decisions.

BASELINE

- â€ Daily support emails: 120
- â€ Average first triage time: 3.5 hours/day
- â€ Misrouting rate: estimated 8â 10%
- â€ Escalation delays: common during peak periods

PILOT DESIGN

- â€ Duration: 30 days
- â€ Scope: one support inbox
- â€ Human review: 100% of AI-prioritized urgent tickets; 25% random sample of routine tickets
- â€ Sensitive data: exclude billing disputes and legal threats from AI processing

METRICS

Metric	Baseline	Target
First triage time	3.5 hr/day	2.0 hr/day
Misrouting rate	8â 10%	below 5%
Urgent issue detection	unknown	measured during pilot
Review burden	unknown	below 45 min/day

KEY RISKS

Risk	Control
Misclassification of urgent issue	Human review and escalation keywords
Privacy exposure	Exclude sensitive categories; approved tool only
Overreliance	AI priority is recommendation, not decision
Customer trust	No external AI-generated response without human review

WORKER FEEDBACK FOCUS

- â€ Does AI reduce repetitive sorting work?
- â€ Does it create extra correction burden?
- â€ Are priority suggestions trustworthy?
- â€ Does it change stress or workload intensity?

DECISION RULE

Scale only if triage time drops, misrouting improves, urgent issues are not missed, and support staff report neutral or positive workload impact.

EXAMPLE: SALES PROPOSAL DRAFTING PILOT

SCENARIO

A small B2B company wants to use AI to prepare first-draft proposal outlines from notes, prior templates, and customer requirements.

AI ROLE

The AI drafts structure and language. Humans own pricing, commitments, technical accuracy, legal terms, and final sending.

BASELINE

- â€ Average proposal drafting time: 4â 6 hours
- â€ Rework from missing requirements: frequent but not tracked
- â€ Proposal quality varies by salesperson

PILOT DESIGN

- â€ Duration: 6 proposal cycles
- â€ Scope: internal draft generation only
- â€ Human review: 100% before external use
- â€ Data rule: no confidential customer data in unapproved tools

METRICS

Metric Baseline Target
--- --- ---
Draft preparation time 4â 6 hr 2â 3 hr

Missing-requirement rework	untracked	track and reduce
Reviewer satisfaction	unknown	4/5 average
Customer-facing errors	0 tolerance	0

KEY RISKS

Risk	Control
Hallucinated product/technical claims	Source-linked requirement checklist
Wrong commercial terms	Pricing/legal fields excluded from AI ownership
Confidentiality breach	Data minimization and approved tools
Deskilling	Use as assistant; train staff on review quality

WORKER FEEDBACK FOCUS

â€¢ Does AI improve draft speed without lowering quality?
â€¢ Are drafts easy to review or subtly misleading?
â€¢ Does the tool help junior staff learn proposal structure?

DECISION RULE

Expand only if draft time improves and human reviewers confirm quality, accuracy, and review burden are acceptable.

EXAMPLE: INTERNAL KNOWLEDGE SEARCH PILOT

SCENARIO

A 75-person organization wants to use AI to search internal policies, SOPs, and project documentation to answer employee questions.

AI ROLE

The AI retrieves and summarizes internal documentation. It must cite source documents. It does not create policy or override managers.

BASELINE

â€¢ Employees ask repeated process questions via chat/email
â€¢ Documentation exists but is hard to find
â€¢ Managers spend time answering repetitive questions

PILOT DESIGN

â€ Duration: 30 days
â€ Scope: non-sensitive SOPs and public internal policies
â€ Human review: weekly sample audit of answers
â€ Data rule: no HR, legal, medical, compensation, or personnel files

METRICS

Metric	Baseline	Target
Repeated manager questions	unknown	reduce by 25%
Answer accuracy	unknown	90%+ sampled accuracy
Source citation rate	n/a	100%
Employee satisfaction	unknown	4/5 average

KEY RISKS

Risk	Control
Outdated answers	document version tracking
Hallucinated policy	require source citation
Unauthorized access	limit indexed documents
False authority	disclose AI is assistant, not policy owner

WORKER FEEDBACK FOCUS

â€ Are answers useful and source-backed?
â€ Does the system reduce interruption burden?
â€ Are employees comfortable asking questions?
â€ Are unclear policies being surfaced for improvement?

DECISION RULE

Scale only if answer accuracy and source citation are high, access controls are reliable, and employees understand the AI is not the final authority.

DISCLAIMER

This toolkit is for practical planning, governance, and risk-management support. It is not legal advice, employment advice, financial advice, compliance certification, or a guarantee that an AI system is safe, lawful, fair, or effective.

High-impact AI uses â including employment, hiring, firing, discipline, pay, scheduling, lending, housing, health, education, legal, safety, insurance, or access to essential services â require qualified review beyond this toolkit.

Organizations remain responsible for their own AI systems, decisions, data practices,

worker engagement, customer communications, and compliance obligations.